



Australian
National
University



In search of ...



Australian
National
University



verifiable...



Australian
National
University



verifiable...



Australian
National
University



verified...



Australian
National
University



on-the-fly...



Australian
National
University



Concurrent Garbage Collectors



Australian
National
University



..on modern processors



Australian
National
University




In search of verified on-the-fly Concurrent Garbage Collection on modern processors

Tony Hosking

Terminology

- The *heap* is a contiguous array of memory words or a set of discontiguous blocks of contiguous words
- A *cell* is a contiguous group of words that may be *allocated* or *free*, or *wasted* or *unusable*

Terminology



- An *object* is a *cell* allocated for use by the application, divided into *fields*
- A *reference* is either a pointer to a heap object or the distinguished value *null*
- A field may contain a *reference* or some other *scalar* non-reference value
- Program *roots* are the local/global variables of a program that hold references

Terminology

- The *mutator* executes application code, which *allocates* new objects and *mutates* references by modifying fields and roots
- Thus, objects can become disconnected from the roots, and become *unreachable*
- The *collector* executes garbage collection code, which discovers unreachable objects and *frees* them, allowing their storage to be *reallocates*

Example

Roots


```
a null
b null
```

Mutator

```
type Cell = record
    value: int;
    next: ref Cell;
end;

var a, b: ref Cell := null;
a := new Cell(0, null);
b := new Cell(1, null);
a.next := b;
b.next := new Cell(2, null);
a := null;
```

Example




Mutator

```
type Cell = record
    value: int;
    next: ref Cell;
end;

var a, b: ref Cell := null;
a := new Cell(0, null);
b := new Cell(1, null);
a.next := b;
b.next := new Cell(2, null);
a := null;
```

Example




Mutator

```
type Cell = record
    value: int;
    next: ref Cell;
end;

var a, b: ref Cell := null;
a := new Cell(0, null);
b := new Cell(1, null);
a.next := b;
b.next := new Cell(2, null);
a := null;
```

Example




Mutator

```
type Cell = record
    value: int;
    next: ref Cell;
end;

var a, b: ref Cell := null;
a := new Cell(0, null);
b := new Cell(1, null);
a.next := b;
b.next := new Cell(2, null);
a := null;
```

Example




Mutator

```
type Cell = record
    value: int;
    next: ref Cell;
end;

var a, b: ref Cell := null;
a := new Cell(0, null);
b := new Cell(1, null);
a.next := b;
b.next := new Cell(2, null);
a := null;
```

Example




Mutator

```
type Cell = record
    value: int;
    next: ref Cell;
end;

var a, b: ref Cell := null;
a := new Cell(0, null);
b := new Cell(1, null);
a.next := b;
b.next := new Cell(2, null);
a := null;
```

Example




Collector: tri-colour marking

```

for each field in Roots
  root := *field
  if root ≠ null && isWhite(root)
    setGrey(root)
    mark()
mark():
  while ∃ source in Grey
    for each field in Pointers(source)
      dest := *field
      if dest ≠ null && isWhite(dest)
        setGrey(dest)
        setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
    mark():
        while ∃ source in Grey
            for each field in Pointers(source)
                dest := *field
                if dest ≠ null && isWhite(dest)
                    setGrey(dest)
                    setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
while ∃ source in Grey
    for each field in Pointers(source)
        dest := *field
        if dest ≠ null && isWhite(dest)
            setGrey(dest)
            setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root != null && isWhite(root)
        setGrey(root)
        mark()
mark():
while  $\exists$  source in Grey
    for each field in Pointers(source)
        dest := *field
        if dest != null && isWhite(dest)
            setGrey(dest)
            setBlack(source)
        
```

Example




Collector: tri-colour marking

```

for each field in Roots
  root := *field
  if root ≠ null && isWhite(root)
    setGrey(root)
    mark()
mark():
  while ∃ source in Grey
    for each field in Pointers(source)
      dest := *field
      if dest ≠ null && isWhite(dest)
        setGrey(dest)
        setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
while ∃ source in Grey
    for each field in Pointers(source)
        dest := *field
        if dest ≠ null && isWhite(dest)
            setGrey(dest)
            setBlack(source)
    
```

Example




Collector: tri-colour marking

```

for each field in Roots
  root := *field
  if root != null && isWhite(root)
    setGrey(root)
    mark()
mark():
  while  $\exists$  source in Grey
    for each field in Pointers(source)
      dest := *field
      if dest != null && isWhite(dest)
        setGrey(dest)
        setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
    while ∃ source in Grey
        for each field in Pointers(source)
            dest := *field
            if dest ≠ null && isWhite(dest)
                setGrey(dest)
                setBlack(source)

```

Example




Collector: tri-colour marking

```

for each field in Roots
  root := *field
  if root ≠ null && isWhite(root)
    setGrey(root)
    mark()
mark():
  while ∃ source in Grey
    for each field in Pointers(source)
      dest := *field
      if dest ≠ null && isWhite(dest)
        setGrey(dest)
        setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
    while ∃ source in Grey
        for each field in Pointers(source)
            dest := *field
            if dest ≠ null && isWhite(dest)
                setGrey(dest)
                setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
  root := *field
  if root ≠ null && isWhite(root)
    setGrey(root)
    mark()
mark():
  while ∃ source in Grey
    for each field in Pointers(source)
      dest := *field
      if dest ≠ null && isWhite(dest)
        setGrey(dest)
        setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
while ∃ source in Grey
    for each field in Pointers(source)
        dest := *field
        if dest ≠ null && isWhite(dest)
            setGrey(dest)
            setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
while ∃ source in Grey
    for each field in Pointers(source)
        dest := *field
        if dest ≠ null && isWhite(dest)
            setGrey(dest)
            setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
    while ∃ source in Grey
        for each field in Pointers(source)
            dest := *field
            if dest ≠ null && isWhite(dest)
                setGrey(dest)
                setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root != null && isWhite(root)
        setGrey(root)
        mark()
mark():
    while  $\exists$  source in Grey
        for each field in Pointers(source)
            dest := *field
            if dest != null && isWhite(dest)
                setGrey(dest)
                setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
    while ∃ source in Grey
        for each field in Pointers(source)
            dest := *field
            if dest ≠ null && isWhite(dest)
                setGrey(dest)
                setBlack(source)
            
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
    while ∃ source in Grey
        for each field in Pointers(source)
            dest := *field
            if dest ≠ null && isWhite(dest)
                setGrey(dest)
            setBlack(source)
        
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
while ∃ source in Grey
    for each field in Pointers(source)
        dest := *field
        if dest ≠ null && isWhite(dest)
            setGrey(dest)
            setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
while ∃ source in Grey
    for each field in Pointers(source)
        dest := *field
        if dest ≠ null && isWhite(dest)
            setGrey(dest)
setBlack(source)
    
```

Example




Collector: tri-colour marking

```

for each field in Roots
  root := *field
  if root ≠ null && isWhite(root)
    setGrey(root)
    mark()
mark():
while ∃ source in Grey
  for each field in Pointers(source)
    dest := *field
    if dest ≠ null && isWhite(dest)
      setGrey(dest)
      setBlack(source)
    
```

Example




Collector: tri-colour marking

```

for each field in Roots
  root := *field
  if root ≠ null && isWhite(root)
    setGrey(root)
    mark()
mark():
while ∃ source in Grey
  for each field in Pointers(source)
    dest := *field
    if dest ≠ null && isWhite(dest)
      setGrey(dest)
      setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
    mark():
        while ∃ source in Grey
            for each field in Pointers(source)
                dest := *field
                if dest ≠ null && isWhite(dest)
                    setGrey(dest)
                    setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
  root := *field
  if root ≠ null && isWhite(root)
    setGrey(root)
    mark()
mark():
  while ∃ source in Grey
    for each field in Pointers(source)
      dest := *field
      if dest ≠ null && isWhite(dest)
        setGrey(dest)
        setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
  root := *field
  if root ≠ null && isWhite(root)
    setGrey(root)
    mark()
mark():
while ∃ source in Grey
  for each field in Pointers(source)
    dest := *field
    if dest ≠ null && isWhite(dest)
      setGrey(dest)
      setBlack(source)
  
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
while ∃ source in Grey
    for each field in Pointers(source)
        dest := *field
        if dest ≠ null && isWhite(dest)
            setGrey(dest)
setBlack(source)
    
```

Example




Collector: tri-colour marking

```

for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
    while ∃ source in Grey
        for each field in Pointers(source)
            dest := *field
            if dest ≠ null && isWhite(dest)
                setGrey(dest)
                setBlack(source)
            
```

Example




Collector: tri-colour marking

```

for each field in Roots
  root := *field
  if root ≠ null && isWhite(root)
    setGrey(root)
    mark()
mark():
  while ∃ source in Grey
    for each field in Pointers(source)
      dest := *field
      if dest ≠ null && isWhite(dest)
        setGrey(dest)
        setBlack(source)
  
```


Example



Collector: sweeping

```
for each cell in Heap
    if isWhite(cell)
        free(cell)
    else
        setWhite(cell)
```


Example



Collector: sweeping

```
for each cell in Heap
  if isWhite(cell)
    free(cell)
  else
    setWhite(cell)
```

Example




Collector: sweeping

```
for each cell in Heap
    if isWhite(cell)
        free(cell)
    else
        setWhite(cell)
```

Example

Roots

a	null
b	




Collector: sweeping

```
for each cell in Heap
    if isWhite(cell)
        free(cell)
    else
        setWhite(cell)
```

Example

Roots

a	null
b	




Collector: sweeping

```
for each cell in Heap  
if isWhite(cell)  
    free(cell)  
else  
    setWhite(cell)
```

Example

Roots

a	null
b	




Collector: sweeping

```
for each cell in Heap
    if isWhite(cell)
        free(cell)
    else
        setWhite(cell)
```

Example

Roots


a	null
b	



Collector: sweeping

```
for each cell in Heap
    if isWhite(cell)
        free(cell)
    else
        setWhite(cell)
```


Example



Collector: sweeping

```
for each cell in Heap  
    if isWhite(cell)  
        free(cell)  
    else  
        setWhite(cell)
```


Example



Collector: sweeping

```
for each cell in Heap
    if isWhite(cell)
        free(cell)
    else
        setWhite(cell)
```


Example



Collector: sweeping

```
for each cell in Heap
    if isWhite(cell)
        free(cell)
    else
        setWhite(cell)
```

Example



Collector: sweeping

```
for each cell in Heap
    if isWhite(cell)
        free(cell)
    else
        setWhite(cell)
```

Correctness Properties

Safety:

The collector never reclaims *live* objects.

More feasibly: the collector never reclaims *reachable* objects.

Liveness:

The collector eventually completes its collection cycle.

Safety Invariants

At the end of each iteration of the mark loop there are no references from black to white objects.

Thus, any *reachable* white object must be reachable from a grey object.

Thus, all *reachable* white objects will eventually be marked.

```
for each field in Roots
    root := *field
    if root ≠ null && isWhite(root)
        setGrey(root)
        mark()
mark():
while ∃ source in Grey
    for each field in Pointers(source)
        dest := *field
        if dest ≠ null && isWhite(dest)
            setGrey(dest)
        setBlack(source)
```

INCREMENTAL AND CONCURRENT COLLECTION

Pauses




The duration of the collection cycle is proportional to the number of marked objects:

- users may notice pauses in interactive applications
- transactions may time out forcing retry
- real-time tasks may miss deadlines


One way to reduce pauses is to use *incremental* or *concurrent collection*.

Uniprocessor incremental collection




Collector *interleaves* with mutator at defined
safe points such as allocation sites

Multiprocessor incremental collection




Collector interleaves with mutator threads

Parallel incremental collection




Collector threads interleave with mutator threads

Mostly-concurrent collection




Mutator threads pause together briefly at the beginning of each collection cycle, then run concurrently with the collector

Mostly-concurrent incremental collection




Mutator threads pause together briefly at the beginning of each collection cycle, then interleave with the collector

On-the-fly collection



No stopping-the-world: mutator threads separately *synchronize* with the collector at the beginning of each collection cycle, then run concurrently with the collector

On-the-fly incremental collection



No stopping-the-world: mutator threads separately *synchronize* with the collector at the beginning of each collection cycle, then *interleave* with the collector

What can go wrong for concurrent GC?


- It's OK for the mutator to modify objects ahead of the waveform (grey or reachable white objects) whose fields are still to be scanned
- It's not OK for the mutator to insert a white pointer into a black object without letting the collector know

The lost object problem

Condition 1: the mutator stores a pointer to a white object into a black object; and


Condition 2: all paths from any grey objects to that white object are destroyed

Example 1: direct




Hiding a reachable white object by dropping a direct link from grey

Example 1: direct




Hiding a reachable white object by dropping a direct link from grey

Example 1: direct




Hiding a reachable white object by dropping a direct link from grey

Example 1: direct




Hiding a reachable white object by dropping a direct link from grey

Example 1: direct




Hiding a reachable white object by dropping a direct link from grey

Example 1: direct




Hiding a reachable white object by dropping a direct link from grey

Example 1: direct




Hiding a reachable white object by dropping a direct link from grey

Example 2: indirect




Hiding a transitively reachable white object by breaking an indirect chain from grey

Example 2: indirect




Hiding a transitively reachable white object by breaking an indirect chain from grey

Example 2: indirect




Hiding a transitively reachable white object by breaking an indirect chain from grey

Example 2: indirect




Hiding a transitively reachable white object by breaking an indirect chain from grey

Example 2: indirect




Hiding a transitively reachable white object by breaking an indirect chain from grey

Example 2: indirect



Hiding a transitively reachable white object by breaking an indirect chain from grey

Example 2: indirect




Hiding a transitively reachable white object by breaking an indirect chain from grey

The strong tricolor invariant

Preventing Condition 1:


Ensure there are no pointers from black objects to white objects

Insertion barrier [Steele]




Retreat the waveform: grey the source
Mutator is allowed to have white roots

Insertion barrier [Steele]




Retreat the waveform: grey the source
Mutator is allowed to have white roots

Insertion barrier [Steele]




Retreat the waveform: grey the source
Mutator is allowed to have white roots

Insertion barrier [Steele]




Retreat the waveform: grey the source
Mutator is allowed to have white roots

Insertion barrier [Steele]




Retreat the waveform: grey the source
Mutator is allowed to have white roots

Insertion barrier [Steele]




Retreat the waveform: grey the source
Mutator is allowed to have white roots

Insertion barrier [Steele]




Retreat the wavefront: grey the source
Mutator is allowed to have white roots

Insertion barrier [Dijkstra]




Advance the wavefront: shade the target
Mutator is allowed to have white roots

Insertion barrier [Dijkstra]




Advance the wavefront: shade the target
Mutator is allowed to have white roots

Insertion barrier [Dijkstra]




Advance the wavefront: shade the target
Mutator is allowed to have white roots

Insertion barrier [Dijkstra]




Advance the waveform: shade the target
Mutator is allowed to have white roots

Insertion barrier [Dijkstra]




Advance the wavefront: shade the target
Mutator is allowed to have white roots

Insertion barrier [Dijkstra]



Advance the waveform: shade the target
Mutator is allowed to have white roots

Insertion barrier [Dijkstra]




Advance the wavefront: shade the target
Mutator is allowed to have white roots

The weak tricolor invariant

Preventing Condition 1:


Every white object pointed to by a black object must be grey protected: that is, reachable from some grey object directly or indirectly via a chain of white objects

Read barrier [Baker]




Advance the waveform: shade the target on load
Mutator is prevented from holding white roots

Read barrier [Baker]




Advance the waveform: shade the target on load
Mutator is prevented from holding white roots

Read barrier [Baker]




Advance the waveform: shade the target on load
Mutator is prevented from holding white roots

Read barrier [Baker]




Advance the wavefront: shade the target on load
Mutator is prevented from holding white roots

Read barrier [Baker]




Advance the wavefront: shade the target on load
Mutator is prevented from holding white roots

Read barrier [Baker]




Advance the wavefront: shade the target on load
Mutator is prevented from holding white roots

Deletion barrier [Abraham & Patel; Yuasa]




Advance the wavefront: shade the deleted target
Weak invariant allows mutator to hold white roots

Deletion barrier [Abraham & Patel; Yuasa]




Advance the wavefront: shade the deleted target
Weak invariant allows mutator to hold white roots

Deletion barrier [Abraham & Patel; Yuasa]




Advance the wavefront: shade the deleted target
Weak invariant allows mutator to hold white roots

Deletion barrier [Abraham & Patel; Yuasa]




Advance the wavefront: shade the deleted target
Weak invariant allows mutator to hold white roots

Deletion barrier [Abraham & Patel; Yuasa]




Advance the waveform: shade the deleted target
Weak invariant allows mutator to hold white roots

Deletion barrier [Abraham & Patel; Yuasa]



Advance the waveform: shade the deleted target
Weak invariant allows mutator to hold white roots

Deletion barrier [Abraham & Patel; Yuasa]




Advance the waveform: shade the deleted target
Weak invariant allows mutator to hold white roots

Termination

Eventually all live objects are black


- Marking only adds to the waveform
- Mutator barriers only add to the waveform
- Allocating black speeds termination
- Deletion barrier avoids need to re-scan mutator roots
 - mutator can safely acquire white roots after snapshot
 - otherwise, must repeat marking from mutator roots until no white roots remain

Allocating black




Deletion barrier is insufficient when allocating black.

Allocating black




Deletion barrier is insufficient when allocating black.

Allocating black




Deletion barrier is insufficient when allocating black.

Allocating black




Deletion barrier is insufficient when allocating black.

Allocating black




Deletion barrier is insufficient when allocating black.

Allocating black




Deletion barrier is insufficient when allocating black.

Allocating black




Deletion barrier is insufficient when allocating black.

Allocating black




Deletion barrier is insufficient when allocating black.

Allocating white




Deletion barrier suffices when allocating white.

Allocating white




Deletion barrier suffices when allocating white.

Allocating white




Deletion barrier suffices when allocating white.

Allocating white




Deletion barrier suffices when allocating white.

Allocating white




Deletion barrier suffices when allocating white.

Allocating white




Deletion barrier suffices when allocating white.

Allocating white




Deletion barrier suffices when allocating white.

Allocating white




Deletion barrier suffices when allocating white.

Allocating grey




Deletion barrier suffices when allocating white/grey.

Allocating grey




Deletion barrier suffices when allocating white/grey.

Allocating grey




Deletion barrier suffices when allocating white/grey.

Allocating grey




Deletion barrier suffices when allocating white/grey.

Allocating grey




Deletion barrier suffices when allocating white/grey.

Allocating grey




Deletion barrier suffices when allocating white/grey.

Allocating grey



Deletion barrier suffices when allocating white/grey.

Allocating grey



Deletion barrier suffices when allocating white/grey.

Safety

Arguments about safety of on-the-fly collectors trade on the weak or strong tricolor invariants enforced by the barriers.

BUT...

Can we really believe these arguments for real?

Can we prove safety for an efficient production algorithm that uses lightweight synchronization on x86 relaxed memory?

Segue from PLDI'10 to PLDI'15 and beyond

PROVING SAFETY

The bones of Schism: CMR [PLDI'10]

Concurrent mark-region GC

- Concurrent
- On-the-fly
- Wait-free constant-time heap access
- Mark-region allocator (linear-time for small objects or no fragmentation)
- Good throughput

CMR data structures

- `o.flag` field indicates if object `o` is marked
- `o.next` field is used to log `o` to a worklist
- `o.len` field object `o` is marked
- `W` is the collector worklist
- `W[t]` is a private worklist for mutator `t`

Mutator primitives

- Read (o, i) : load slot i from object o
- Write (o, i, v) : store v into slot i of o
- Alloc (n, v) : return a reference to a new object having n fields initialized to v
- Collector-mutator *handshakes* in between,
never in the middle

CMR Write

$\text{Write}(o, i, v) :$

$w := o[i]$

$\text{mark}(w, W[\text{self}])$

$\text{mark}(v, W[\text{self}])$

$o[i] := v$

CMR Alloc

Alloc(n, v) :

- o := allocRaw(n)

- o.flag := fA

- o.next := null

- o.len := n

for each l in Refs(o)

- *l := v

return o

CMR Collector

loop

```

phase ← Idle
handshake t in T nop
 $f_M \leftarrow !f_M$ 
handshake t in T nop
phase ← Init
handshake t in T nop
phase ← Mark
 $f_A \leftarrow f_M$ 
handshake t in T nop
handshake t in T
    for each l in Roots(t)
        mark(*l, W[t])
    atomic transfer(W[t], W)

```

```

while W.head ≠ null
    while W.head ≠ null
        s ← dequeue(W)
        shadeBlock(s, s.len)
        for each l in Refs(s)
            mark(*l, W)
        handshake t in T
            atomic transfer(W[t], W)
    phase ← Sweep
    sweepBlocksAndLines()

```

CMR Collector

loop

```

phase ← Idle
handshake  $t \in T$  nop
 $f_M \leftarrow !f_M$ 
handshake  $t \in T$  nop
phase ← Init
handshake  $t \in T$  nop
phase ← Mark
 $f_A \leftarrow f_M$ 
handshake  $t \in T$  nop
handshake  $t \in T$ 
    for each  $l \in \text{Roots}(t)$ 
        mark( $*l, W[t]$ )
    atomic transfer( $W[t], W$ )

```

idle: no marking

```

while  $W.\text{head} \neq \text{null}$ 
    while  $W.\text{head} \neq \text{null}$ 
         $s \leftarrow \text{dequeue}(W)$ 
        shadeBlock( $s, s.\text{len}$ )
        for each  $l \in \text{Refs}(s)$ 
            mark( $*l, W$ )
    handshake  $t \in T$ 
    atomic transfer( $W[t], W$ )
phase ← Sweep
sweepBlocksAndLines()

```

CMR Collector

loop

```

phase ← Idle
handshake t in T nop
 $f_M \leftarrow !f_M$ 
handshake t in T nop
phase ← Init
handshake t in T nop
phase ← Mark
 $f_A \leftarrow f_M$ 
handshake t in T nop
handshake t in T
    for each l in Roots(t)
        mark(*l, W[t])
    atomic transfer(W[t], W)

```

```

while W.head ≠ null
    while W.head ≠ null
        s ← dequeue(W)
        shadeBlock(s, s.len)
        for each l in Refs(s)
            mark(*l, W)
        handshake t in T
            atomic transfer(W[t], W)
    phase ← Sweep
    sweepBlocksAndLines()

```

CMR Collector

loop

```

phase ← Idle
handshake  $t \in T$  nop
 $f_M \leftarrow !f_M$ 
handshake  $t \in T$  nop
phase ← Init
handshake  $t \in T$  nop
phase ← Mark
 $f_A \leftarrow f_M$ 
handshake  $t \in T$  nop
handshake  $t \in T$ 
    for each  $l \in \text{Roots}(t)$ 
        mark( $*l, W[t]$ )
    atomic transfer( $W[t], W$ )

```

while $W.\text{head} \neq \text{null}$

while $W.\text{head} \neq \text{null}$

$s \leftarrow \text{dequeue}(W)$

shadeBlock($s, s.\text{len}$)

for each $l \in \text{Refs}(s)$

 mark($*l, W$)

handshake $t \in T$

atomic transfer($W[t], W$)

phase ← Sweep

sweepBlocksAndLines()

Init: Store marks

CMR Collector

loop

```

phase ← Idle
handshake t in T nop
 $f_M \leftarrow !f_M$ 
handshake t in T nop
phase ← Init
handshake t in T nop
phase ← Mark
 $f_A \leftarrow f_M$ 
handshake t in T nop
handshake t in T
    for each l in Roots(t)
        mark(*l, W[t])
    atomic transfer(W[t], W)

```

```

while W.head ≠ null
    while W.head ≠ null
        s ← dequeue(W)
        shadeBlock(s, s.len)
        for each l in Refs(s)
            mark(*l, W)
        handshake t in T
            atomic transfer(W[t], W)
    phase ← Sweep
    sweepBlocksAndLines()

```

CMR Collector

loop

```

phase ← Idle
handshake  $t \in T$  nop
 $f_M \leftarrow !f_M$ 
handshake  $t \in T$  nop
phase ← Init
handshake  $t \in T$  nop
phase ← Mark
 $f_A \leftarrow f_M$ 
handshake  $t \in T$  nop
handshake  $t \in T$ 
    for each  $l \in \text{Roots}(t)$ 
        mark(* $l$ ,  $W[t]$ )
    atomic transfer( $W[t]$ ,  $W$ )

```

```

while  $W.\text{head} \neq \text{null}$ 
    while  $W.\text{head} \neq \text{null}$ 
         $s \leftarrow \text{dequeue}(W)$ 
        shadeBlock( $s, s.\text{len}$ )
        for each  $l \in \text{Refs}(s)$ 
            mark(* $l$ ,  $W$ )
        handshake  $t \in T$ 
        atomic transfer( $W[t]$ ,
 $W$ )
    phase ← Sweep
    sweepBlocksAndLines()

```

Mark: Store &
Alloc mark

CMR Collector

loop

```

phase ← Idle
handshake t in T nop
 $f_M \leftarrow !f_M$ 
handshake t in T nop
phase ← Init
handshake t in T nop
phase ← Mark
 $f_A \leftarrow f_M$ 
handshake t in T nop
handshake t in T
    for each l in Roots(t)
        mark(*l, W[t])
    atomic transfer(W[t], W)

```

```

while W.head ≠ null
    while W.head ≠ null
        s ← dequeue(W)
        shadeBlock(s, s.len)
        for each l in Refs(s)
            mark(*l, W)
        handshake t in T
            atomic transfer(W[t], W)
    phase ← Sweep
    sweepBlocksAndLines()

```

CMR Collector

loop

```

phase ← Idle
handshake t in T nop
 $f_M \leftarrow !f_M$ 
handshake t in T nop
phase ← Init
handshake t in T nop
phase ← Mark
 $f_A \leftarrow f_M$ 
handshake t in T nop
handshake t in T
    for each l in Roots(t)
        mark(*l, W[t])
    atomic transfer(W[t], W)

```

```

while W.head ≠ null
    while W.head ≠ null
        s ← dequeue(W)
        shadeBlock(s, s.len)
        for each l in Refs(s)
            Sweep:Alloc marks mark(*l, W)
            handshake t in T
            atomic transfer(W[t], W)
    phase ← Sweep
    sweepBlocksAndLines()

```

CMR Collector

loop

```

phase ← Idle
handshake t in T nop
 $f_M \leftarrow !f_M$ 
handshake t in T nop
phase ← Init
handshake t in T nop
phase ← Mark
 $f_A \leftarrow f_M$ 
handshake t in T nop
handshake t in T
    for each l in Roots(t)
        mark(*l, W[t])
    atomic transfer(W[t], W)

```

```

while W.head ≠ null
    while W.head ≠ null
        s ← dequeue(W)
        shadeBlock(s, s.len)
        for each l in Refs(s)
            mark(*l, W)
        handshake t in T
            atomic transfer(W[t], W)
    phase ← Sweep
    sweepBlocksAndLines()

```

CMR Mark

mark(\circ , W) :

if $\circ.flag \neq fM$

if phase \neq Idle

if CAS(& $\circ.flag$, ! $fM \rightarrow fM$) $=$! fM

$\circ.next := W.head$

if $W.tail = null$

$W.tail := \circ$

$W.head := \circ$

CMR Transfer

```
transfer(W1, W2) :  
    if W1.head ≠ null  
        W1.tail.next := W2.head  
        W2.head := W1.head  
        W1.head := null  
        W1.tail := null
```

CMR Dequeue

dequeue (W) :

$o \leftarrow W.\text{head}$

if $o = \text{null}$

return null

$W.\text{head} \leftarrow o.\text{next}$

if $W.\text{tail} = o$

$W.\text{tail} = \text{null}$

return o



Relaxing Safely: Verified On-the-Fly Garbage Collection for x86-TSO

Peter Gammie¹ Tony Hosking² Kai Engelhardt³

¹ex NICTA Australia


²Purdue University while on leave at NICTA

³UNSW and NICTA Australia

Tony's Goal

Verify a highly concurrent, on-the-fly garbage collector wrt a non-sequentially consistent memory model

Our headline result



For $Sys \simeq TSO$, there is always an object at every reference reachable from a mutator root

Main challenges


- ▶ the system is **highly concurrent**
- ▶ memory is not **sequentially consistent**
- ▶ mutators may not be **data-race free**
- ▶ **convince**  of our proof technique/invariants

The Schism Garbage Collector

as described at PLDI'10

Concurrent, on-the-fly, mark-sweep

- ▶ Does **not stop the world**
 - ▶ wait-free for the mutators up to allocation
- ▶ Takes a **snapshot**
- ▶ Races with the mutators to mark reachable objects
- ▶ Objects are built from cache-line-sized chunks
 - ▶ representation bakes in GC space overhead
 - ⇒ fragmentation tolerant without relocation



Claim: has **predictable real-time performance**

Folklore: GC requires mutator cooperation


Mutators use **write barriers** unless the GC is idle

- ▶ the **insertion barrier** prevents mutators from hiding references from the collector
- ▶ the **deletion barrier** accelerates termination

Objects are allocated...


- ▶ **white** when GC is **idle**
 - ▶ minimizes **floating garbage**
- ▶ **black** when GC is **not idle**
 - ▶ accelerates termination

Grey objects track pending GC work



Total Store Order (x86-TSO)

x86-TSO has **one write queue per core** and **consults these queues for memory reads**



Some instructions (CAS) **are an atomic sequence memory operations**, so there is also a **bus lock**

x86-TSO is **not sequentially consistent**
... there can be **data races**

- ▶ The model has been **validated against recent x86 hardware** by Peter Sewell et al.

Data races

Starting in a state where $x = y = 0$:

P_0	P_1
$x \leftarrow 1$	$y \leftarrow 1$
$r_1 \leftarrow y$	$r_2 \leftarrow x$

Can observe $r_1 = r_2 = 0$
a data race

Data races

Starting in a state where $x = y = 0$:

P_0	P_1
$x \leftarrow 1$	$y \leftarrow 1$
MFENCE	MFENCE
$r_1 \leftarrow y$	$r_2 \leftarrow x$

MFENCE waits until the core's TSO buffer is empty

Now **cannot** observe $r_1 = r_2 = 0$

We have sequential consistency!

... with significantly degraded performance

The promise of many reduction theorems


If there are **no data races**, then we can use **classical techniques** for sequentially consistent memory

Problem: our mutators need not be **DRF**

Solution: resort to general techniques for concurrency




So much for the model... what about the assertion?




For $\text{Sys} \simeq \text{TSO}$, there is always an object at every reference reachable from a mutator root

So much for the model... what about the assertion?



For $\text{Sys} \simeq \text{TSO}$, there is always *an object at every reference* reachable from a mutator root


So much for the model... what about the assertion?




For $Sys \simeq TSO$, there is always an object at every reference *reachable from a mutator root*

What does TSO do to reachability?


TSO and reachability, breathlessly



TSO and reachability, breathlessly




TSO and reachability, breathlessly



The pending write is the only witness to B's reachability

TSO and reachability, breathlessly




... but the **insertion barrier has already greyed B!**


We treat grey and TSO refs as **extra mutator roots**

... and **all paths go via the heap**

So much for the result... how did we express all that?



For $Sys \simeq TS0$, there is always an object at every reference reachable from a mutator root



"You can't always write a chord ugly enough to say what you want to say, so sometimes you have to rely on a giraffe filled with whipped cream."

The framework: CIMP

- ▶ Simple imperative language (IMP) + synchronous message passing
- ▶ Flat, top-level parallel composition.
- ▶ Each process has local control and data states

There is **no shared global state**

- ▶ Amenable to **state-based invariant reasoning**

Enables **separation** of programs, invariants and proofs

Constructing the GC invariants

Tony was confident that the GC satisfied the **standard mark/sweep tricolour invariants**

. . . but before we made it to those lofty realms, we first:

- ▶ Defined a system-wide “**program counter**” that encodes the GC’s phase structure
- ▶ **Exploited DRF** where we could, and sliced the program wrt non-DRF variables when we had to
- ▶ Developed **fine-grained assertions** around the **object marking operations**

. . . in higher-order logic



Isabelle’s parallelism for low-latency invariant search!

Are you advocating this for arbitrary programs?

No

This was about discovering invariants for a particular non-trivial program on x86-TSO

Why didn't you develop or use...

- ▶ a suitable reduction theorem?
- ▶ something compositional, like rely/guarantee?
- ▶ separation logic and ownership?
- ▶ shared variables and Owicki/Gries?
- ▶ communication-closed layers?
- ▶ a more abstract model and refinement?
- ▶ some other formalism such as I/O automata?
- ▶ proof outlines?
- ▶ a modelchecker?
- ▶ ... your favoured technique?

Why didn't you develop or use...

- ▶ a suitable reduction theorem?
- ▶ something compositional, like rely/guarantee?
- ▶ separation logic and ownership?
- ▶ shared variables and Owicki/Gries?
- ▶ communication-closed layers?
- ▶ a more abstract model and refinement?
- ▶ some other formalism such as I/O automata?
- ▶ proof outlines?
- ▶ a modelchecker?
- ▶ ... your favoured technique?

Largely because (almost) none of these would have eased invariant discovery.

OK, so where did you cheat?

Minor:

- ▶ grey is not subject to TSO
 - ▶ underlying memory blocks are not modelled
 - ▶ ragged safepoints are abstract
 - ▶ each mutator + GC runs on its own core
- ⇒ We leave these to an atomicity refinement technique.

Major:

- ▶ alloc and free are global and atomic.

Remaining and future work

- ▶ A defensible treatment of allocation
- ▶ Connection with something more executable
 - ▶ x86 instruction semantics + refinement
- ▶ Recast in a compositional framework
 - ▶ extra invariants: the heap is not mutilated, . . .

A plausible liveness result may prove elusive.

Come and talk to us if any of that interests you!

Concluding remarks

- ▶ Reasoning about data-racy programs on x86-TSO is (sometimes) not too bad... but that might be because the collector is carefully constructed
- ▶ To scale one certainly wishes to exploit the general absence of data races
- ▶ The x86-TSO model is useful for thinking about correctness but not performance
 - ▶ How do we provide analytic WCET bounds for this putatively real-time collector?
- ▶ ARM/POWER are more complex